

Setup Guide

This guide was written as a support for the different Kubernetes courses that I'm teaching. This guide does not intend to show all possible solutions but covers a few options to use Kubernetes. Other options do exist, but I cannot support them in my courses. For questions: mail@sandervanvugt.nl

Setting up Minikube on a Linux virtual machine

Any virtualization platform will do, but these steps are for VMware Fusion/Workstation or VirtualBox. Installation on MacOS or Windows is also possible but may give additional problems and is for that reason not recommended.

Note: In nov. 2020 I have experienced problems running nested virtualization in VMware Fusion 12 Pro on Mac OS Big Sur. If you have this configuration, you might want to go for the procedure for Setting up AiO Kubernetes described next.

1. Make sure you have either a fully installed Fedora 33 or Ubuntu 20.04 virtual machine
 - a. 8 GB RAM recommended, 4 GB RAM minimal
 - b. CPU's in the virtual machine have nested virtualization enabled
 - i. On VMware: select VM > Processors and Memory > Advanced Options and select both "Enable Hypervisor applications in this virtual machine" and "Enable IOMMU..."
 - c. 40 GB disk space recommended, 20 GB minimal
2. Install git, vim and bash completion
 - a. On Ubuntu: **sudo apt install git vim bash-completion**
 - b. On Fedora: **sudo dnf install git vim bash-completino**
3. As ordinary user with sudo privileges, clone the course Git repository
 - a. **git clone https://github.com/sandervanvugt/kubernetes** for Kubernetes in 4 Hours
 - b. **git clone https://github.com/sandervanvugt/ckad** for CKAD
 - c. **git clone https://github.com/sandervanvugt/microservices** for Microservices
4. Change into the cloned git repo and run the **kube-setup.sh** script
 - a. **cd kubernetes**
 - b. **./kube-setup**
5. At the end of the script, reboot your virtual machine and run the following:
 - a. **minikube start --memory 4096 --vm-driver=kvm2**
6. Once Minikube has started successfully, you'll see a message that it has started. Type **kubectl get all** to verify the minikube cluster is up and running

Setting up AiO-Kubernetes

This method is less common, but offers the advantage that no nested virtualization is required.

1. Install a Centos 7 (NOT 8) or Ubuntu 20.04 Server virtual machine; minimal installation will do
 - a. 20 GB disk space
 - b. 4 GB RAM recommended (2 GB minimal)
 - c. 2 CPU's
 - d. No Swap
2. Install some packages
 - a. On Ubuntu: **apt install vim git bash-completion**
 - b. On CentOS: **yum install git vim bash-comletion**
3. As ordinary user with sudo privileges, clone the course Git repository
 - a. **git clone https://github.com/sandervanvugt/kubernetes** for Kubernetes in 4 Hours
 - b. **git clone https://github.com/sandervanvugt/ckad** for CKAD
 - c. **git clone https://github.com/sandervanvugt/microservices** for Microservices
4. Run the setup scripts:
 - a. **cd /kubernetes** (or whichever GitHub repository you have clones)
 - b. **./setup-docker.sh**
 - c. **./setup-kubetools.sh**
5. In a root shell, install a Kubernetes master node
 - a. **kubeadm init**
6. In a user shell, set up the kubectl client:
 - a. **mkdir -p \$HOME/.kube**
 - b. **sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config**
 - c. **sudo chown \$(id -un):\$(id -un) .kube/config**
7. In a user shell, set up the network agent (all on one line)
 - a. **kubectl apply -f https://cloud.weave.works/k8s/net?k8s-version=\$(kubectl version | base64 | tr -d '\n')**
8. Use **kubectl get nodes** to find the current node name. Replace **kubeaio.example.com** in the next command with the node name you have found
9. Open the kubectl editor to chage node properties:
 - a. **kubectl edit node kubeaio.example.com**
10. Look up the following lines and delete all of them:

taints:

- effect: NoSchedule

key: node-role.kubernetes.io/master
11. Type **kubectl get all** to verify the cluster works.

Using Hosted Kubernetes in GCE

Kubernetes is supported by many cloud providers. As I cannot be all-inclusive, but most of all, as Google has donated Kubernetes to the world, I like to do something back and explain the procedure on GCE only.

1. From a browser, go to <https://cloud.google.com> and click **Sign in**.

2. Click **Go to console**
3. Select **Compute > Kubernetes Engine > Clusters**
4. Click **Create Cluster**
5. Click **My first cluster**
6. Click **Create Now** and wait a few minutes - typically not more than 5.
7. Once the cluster appears, click **Connect**
8. You'll see a command **gcloud container ...**. Click **Run in Cloud Shell** to run it. This will deploy a cloud shell machine, which may take one or two minutes.
9. Press Enter to run the command, and click **Authorize** to authorize the cloud shell
10. Type **kubectl get all** to check that the Kubernetes cluster works.